

# Lecture 15: Bagging, Random Forests

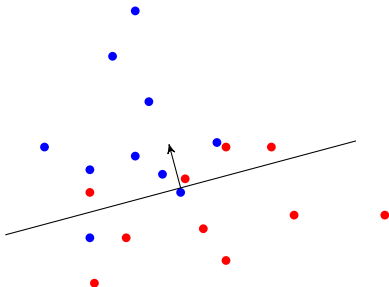
Reading: Sections 9.2, 15.2, 15.3

GU4241/GR5241 Statistical Machine Learning

Linxi Liu  
April 5, 2019


# Ensembles

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



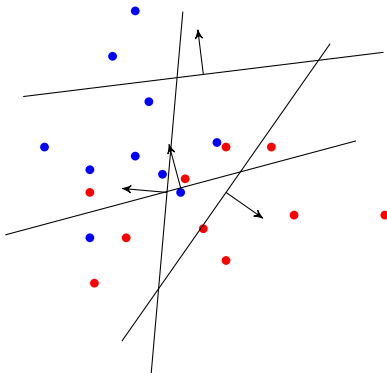
# Ensembles

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



# Ensembles

A *randomly* chosen hyperplane classifier has an *expected* error of 0.5 (i.e. 50%).



- ▶ Many random hyperplanes combined by majority vote: Still 0.5.
- ▶ A single classifier slightly better than random:  $0.5 + \epsilon$ .
- ▶ What if we use  $m$  such classifiers and take a majority vote?

# Voting

## Decision by majority vote

- ▶  $m$  individuals (or classifiers) take a vote.  $m$  is an odd number.
- ▶ They decide between two choices; one is correct, one is wrong.
- ▶ After everyone has voted, a decision is made by simple majority.

**Note:** For two-class classifiers  $f_1, \dots, f_m$  (with output  $\pm 1$ ):

$$\text{majority vote} = \text{sgn}\left(\sum_{j=1}^m f_j\right)$$

## Assumptions

Before we discuss ensembles, we try to convince ourselves that voting can be beneficial. We make some simplifying assumptions:

- ▶ Each individual makes the right choice with probability  $p \in [0, 1]$ .
- ▶ The votes are *independent*, i.e. stochastically independent when regarded as random outcomes.

# Does the Majority Make the Right Choice?


## Condorcet's rule

If the individual votes are independent, the answer is

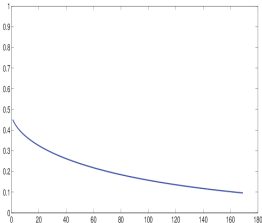
$$\Pr\{\text{majority makes correct decision}\} = \sum_{j=\frac{m+1}{2}}^m \frac{m!}{j!(m-j)!} p^j (1-p)^{m-j}$$

This formula is known as **Condorcet's jury theorem**.

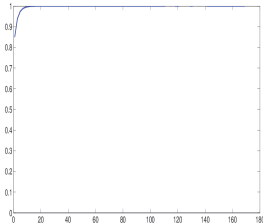
Probability as function of the number of votes



$p = 0.55$



$p = 0.45$



$p = 0.85$

# Ensemble Methods

## Terminology

- ▶ An **ensemble method** makes a prediction by combining the predictions of many classifiers into a single vote.
- ▶ The individual classifiers are usually required to perform only slightly better than random. For two classes, this means slightly more than 50% of the data are classified correctly. Such a classifier is called a **weak learner**.

## Strategy

- ▶ We have seen above that if the weak learners are random and independent, the prediction accuracy of the majority vote will increase with the number of weak learners.
- ▶ Since the weak learners all have to be trained on the training data, producing random, independent weak learners is difficult.
- ▶ Different ensemble methods (e.g. Boosting, Bagging, etc) use different strategies to train and combine weak learners that behave relatively independently.

# Methods We Will Discuss

## Boosting

- ▶ After training each weak learner, data is modified using weights.
- ▶ Deterministic algorithm.

## Bagging

Each weak learner is trained on a random subset of the data.

## Random forests

- ▶ Bagging with decision trees as weak learners.
- ▶ Uses an additional step to remove dimensions in  $\mathbb{R}^d$  that carry little information.



## Bagging = Bootstrap Aggregation

- ▶ Replicate the dataset by sampling with replacement.
- ▶ We apply a learning method to each bootstrap replicate, to produce predictions  $\hat{f}^{(1)}, \dots, \hat{f}^{(B)}$ .
- ▶ In Chapter 5, we were interested in the variability of these predictions:

$$\text{SE}(\hat{f}(x)) \approx \text{SD}(\hat{f}^{(1)}(x), \dots, \hat{f}^{(B)}(x)).$$

- ▶ Now, we will use the average of these predictions as an estimator with reduced variance:

$$\hat{f}^{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}(x)$$

## Bagging decision trees

- ▶ Replicate the dataset by sampling with replacement.
- ▶ Fit a decision tree to each bootstrap replicate (growing the tree, and pruning).
- ▶ **Regression:** To make a prediction for an input point  $x$ , average the predictions of all the trees:

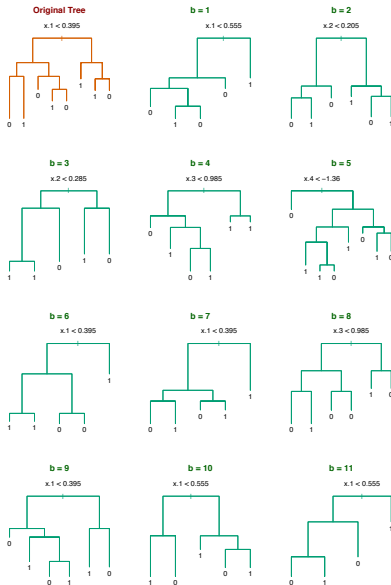
$$\hat{f}^{\text{bag}}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{(b)}(x)$$

- ▶ **Classification:** To make a prediction for an input point  $x_0$ , take the majority vote from the set of predictions:

$$\hat{y}_0^{(1)}, \dots, \hat{y}_0^{(B)}.$$

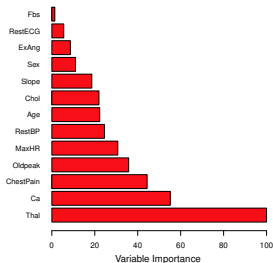
## Example: Bagging decision trees

- ▶ Two classes, each with Gaussian distribution in  $\mathbb{R}^5$ .
- ▶ Note the variance between bootstrapped trees.



## Bagging decision trees

- ▶ **Disadvantage:** Every time we fit a decision tree to a Bootstrap sample, we get a different tree  $T^b$ .  
→ Loss of interpretability
- ▶ For each predictor, add up the total amount by which the RSS (or Gini index) decreases every time we use the predictor in  $T^b$ .
- ▶ Average this total over each Bootstrap estimate  $T^1, \dots, T^B$ .



## How Often Do We See Each Sample in Bootstrap?

Sample  $\{\mathbf{x}_1, \dots, \mathbf{x}_n\}$ , bootstrap resamples  $\mathcal{B}_1, \dots, \mathcal{B}_B$ .

In how many sets does a given  $\mathbf{x}_i$  occur?

Probability for  $\mathbf{x}_i$  *not* to occur in  $n$  draws:

$$\Pr\{\mathbf{x}_i \notin \mathcal{B}_b\} = \left(1 - \frac{1}{n}\right)^n$$

For large  $n$ :

$$\lim_{n \rightarrow \infty} \left(1 - \frac{1}{n}\right)^n = \frac{1}{e} \approx 0.3679$$

- ▶ Asymptotically, any  $\mathbf{x}_i$  will appear in  $\sim 63\%$  of the bootstrap resamples.
- ▶ Multiple occurrences possible.


How often is  $\mathbf{x}_i$  expected to occur?

The *expected* number of occurrences of each  $\mathbf{x}_i$  is  $B$ .

## Out-of-bag (OOB) error

- ▶ To estimate the test error of a bagging estimate, we could use cross-validation.
- ▶ Each time we draw a Bootstrap sample, we only use 63% of the observations.
- ▶ **Idea:** use the rest of the observations as a test set.
- ▶ **OOB error:**
  - ▶ For each sample  $x_i$ , find the prediction  $\hat{y}_i^b$  for all bootstrap samples  $b$  which do not contain  $x_i$ . There should be around  $0.37B$  of them. Average these predictions to obtain  $\hat{y}_i^{\text{oob}}$ .
  - ▶ Compute the error  $(y_i - \hat{y}_i^{\text{oob}})^2$ .
  - ▶ Average the errors over all observations  $i = 1, \dots, n$ .
- ▶ For  $B$  large, OOB error is virtually equivalent to LOOCV.

## Out-of-bag (OOB) error



The test error decreases as we increase  $B$   
(dashed line is the error for a plain decision tree).

# Random Forests

Bagging has a problem:


→ The trees produced by different Bootstrap samples can be very similar.

## Random Forests:


- ▶ We fit a decision tree to different Bootstrap samples.
- ▶ When growing the tree, we select a random sample of  $m < p$  predictors to consider in each step.
- ▶ This will lead to very different (or “uncorrelated”) trees from each sample.
- ▶ Finally, average the prediction of each tree.



# Random Forests vs. Bagging



## Random Forests, choosing $m$



The optimal  $m$  is usually around  $\sqrt{p}$ ,  
but this can be used as a tuning parameter.